

Decision Trees for Iris Flower Classification - Project Overview

Project Objectives

This notebook presents a comprehensive introduction to **Decision Tree classification** using the scikit-learn library and the Iris dataset. The primary objectives are:

1. **Understand Decision Tree Algorithm:** Learn how decision trees recursively split data based on feature values to build interpretable classification models
2. **Model Training:** Train decision tree classifiers on the Iris dataset with various hyperparameter configurations
3. **Tree Visualization:** Visually represent the learned tree structure to understand decision rules and splits
4. **Feature Importance Analysis:** Identify which morphological features (sepal/petal measurements) are most discriminative
5. **Model Evaluation:** Assess performance using accuracy metrics, confusion matrices, and classification reports
6. **Hyperparameter Tuning:** Explore the impact of tree depth, minimum samples split, and other parameters on model complexity and performance

Decision Tree Fundamentals

What Are Decision Trees?

Decision trees are supervised learning models that make predictions by recursively partitioning the feature space into progressively smaller regions. The tree structure consists of:

- **Root node:** The initial split using the most informative feature
- **Internal nodes:** Intermediate decision points representing feature comparisons (e.g., "petal length \leq 2.45 cm?")
- **Branches:** Edges representing outcome paths (yes/no, left/right)
- **Leaf nodes:** Terminal nodes representing final class predictions (setosa/versicolor/virginica)

Each split is chosen to maximize **information gain** (reduction in entropy) or minimize **Gini impurity**, creating the purest possible child nodes at each step.

Why Decision Trees for Iris?

Decision trees are particularly well-suited for the Iris classification problem because:

- **Interpretability:** Results in human-readable rules (e.g., "if petal width > 1.8 → virginica")
- **No preprocessing required:** Handles different feature scales automatically without normalization
- **Non-linear boundaries:** Captures complex decision boundaries between species
- **Feature selection:** Automatically identifies the most important features
- **Efficiency:** Fast training and prediction even on modern hardware
- **Clear class separability:** Iris data has natural hierarchical structure that trees exploit well

Dataset and Methodology

Data Description

The analysis uses the complete Iris dataset consisting of:

- **150 samples:** 50 from each of three iris species
- **4 features:** Sepal length, sepal width, petal length, petal width (all in cm)
- **3 classes:** Setosa, Versicolor, Virginica
- **Balanced distribution:** Equal representation ensures unbiased tree learning

Implementation Workflow

Phase 1: Data Preparation

The notebook loads the Iris dataset from scikit-learn, extracts features (X) and target labels (y), and prepares data in the format required for training. The dataset is examined to verify structure and dimensions before model training.

Phase 2: Train-Test Split

The data is typically divided into:

- **Training set (70-80%):** Used to learn tree structure and splits
- **Test set (20-30%):** Reserved for unbiased model evaluation

This separation prevents overfitting and provides realistic performance estimates on unseen data.

Phase 3: Decision Tree Training

A `DecisionTreeClassifier` is instantiated from scikit-learn with specified hyperparameters and trained on the training data using the `.fit()` method. The algorithm automatically learns the optimal tree structure by:

- Selecting the feature providing maximum information gain at each node
- Creating binary splits that recursively partition the data
- Stopping when pure leaf nodes are reached or other stopping criteria are met

Phase 4: Tree Visualization

The learned tree structure is visualized using `plot_tree()` or exported in various formats, revealing:

- Decision rules at each split
- Feature and threshold values used for splitting
- Number of samples in each node
- Class distribution and purity at each node
- Actual predictions at leaf nodes

Phase 5: Feature Importance Extraction

The `feature_importances_` attribute quantifies each feature's contribution to classification decisions:

- **Petal features** typically dominate (values 0.6-0.9)
- **Sepal features** have lower importance (values 0.1-0.4)
- Importance is calculated based on information gain across all splits

Phase 6: Model Evaluation

Comprehensive evaluation using multiple metrics:

- **Accuracy**: Percentage of correct predictions on test set
- **Precision and Recall**: Per-class performance metrics
- **Confusion Matrix**: Shows misclassification patterns between species
- **Classification Report**: Detailed breakdown of precision, recall, and F1-scores

Hyperparameter Exploration

The notebook explores how different hyperparameters affect tree behavior:

max_depth: Controls tree depth

- Shallow trees (depth 3-4) prevent overfitting but may underfit
- Deep trees (depth 10+) risk memorizing training data
- Optimal depth typically 3-5 for Iris dataset

min_samples_split: Minimum samples required to split a node

- Higher values (10-20) create simpler, more general trees
- Lower values (2-5) allow more detailed splits

min_samples_leaf: Minimum samples required at leaf nodes

- Prevents creation of very small, pure subsets
- Higher values generalize better

Expected Results and Findings

Tree Structure for Iris

A typical learned tree exhibits this decision path:

```
text
Root: petal_length ≤ 2.45 cm?
└─ YES (Left): → Setosa (pure node, 100%)
└─ NO (Right): petal_width ≤ 1.75 cm?
    └─ YES (Left): → Versicolor (mostly pure)
    └─ NO (Right): → Virginica (mostly pure)
```

Performance Metrics

Expected performance levels:

- **Overall Accuracy**: 95-100%
- **Setosa Classification**: 100% (perfectly separable)
- **Versicolor/Virginica**: 90-98% (some overlap in measurements)
- **Tree Depth**: Typically 3-5 levels for good generalization
- **Number of Leaves**: Usually 4-8 terminal nodes

Feature Importance Rankings

Typical importance distribution:

1. **Petal length**: 40-50% importance (best discriminator)
2. **Petal width**: 40-50% importance (equally important)
3. **Sepal length**: 5-15% importance (weak discriminator)
4. **Sepal width**: <5% importance (nearly irrelevant)

This aligns with prior findings from exploratory analysis showing petal features as primary separators.

Key Insights and Learning Outcomes

Algorithmic Insights

- Trees learn hierarchical decision rules that humans can interpret and validate
- First splits use most discriminative features (petal length and width)
- Subsequent splits fine-tune classification between overlapping classes
- Tree structure evolves differently with various hyperparameter choices

Domain Insights

- Iris species separation is achievable with just 2-3 petal measurements
- Sepal measurements provide minimal additional information
- Different depth/complexity levels achieve similar accuracy, with shallower trees preferred for generalization

Practical Lessons

- Decision trees require no feature scaling, unlike many other algorithms
- Tree depth control is critical to balance underfitting and overfitting
- Visualization of learned trees aids in understanding model decisions and debugging
- Feature importance analysis validates domain knowledge and data characteristics

Advantages and Disadvantages of Decision Trees

Advantages

- **Interpretability:** Easy to understand and explain decisions to non-technical stakeholders
- **No preprocessing:** Works with raw features; no normalization required
- **Automatic feature selection:** Identifies important features naturally
- **Fast predictions:** Simple to deploy and execute
- **Handles non-linear relationships:** Can capture complex patterns
- **Works with mixed data types:** Handles both numerical and categorical features

Disadvantages

- **Overfitting tendency:** Deep trees memorize training noise
- **Unstability:** Small data changes cause large tree structure changes
- **Biased toward high-cardinality features:** Prefers features with many split opportunities
- **Poor on linearly separable data:** Axis-aligned splits inefficient for tilted boundaries
- **Limited extrapolation:** Cannot predict beyond observed feature ranges

Extensions and Related Topics

This notebook serves as a foundation for advanced ensemble methods:

- **Random Forests:** Combines multiple trees to reduce overfitting and improve robustness
- **Gradient Boosting:** Sequentially builds trees to correct previous errors
- **Bagging:** Bootstrap samples create diverse trees that vote on predictions

All these methods inherit decision trees' interpretability while improving predictive performance.

Conclusion

This project demonstrates that decision trees provide an effective and interpretable approach to iris flower classification. By systematically training, visualizing, and evaluating tree models with various hyperparameters, practitioners gain both algorithmic understanding and practical experience with a fundamental machine learning technique. The analysis validates that tree-based methods successfully separate iris species based on morphological features, achieving near-perfect classification performance while maintaining model interpretability—a critical advantage in domains requiring explainable predictions.